

# TCP/IP

EN.600.444/644

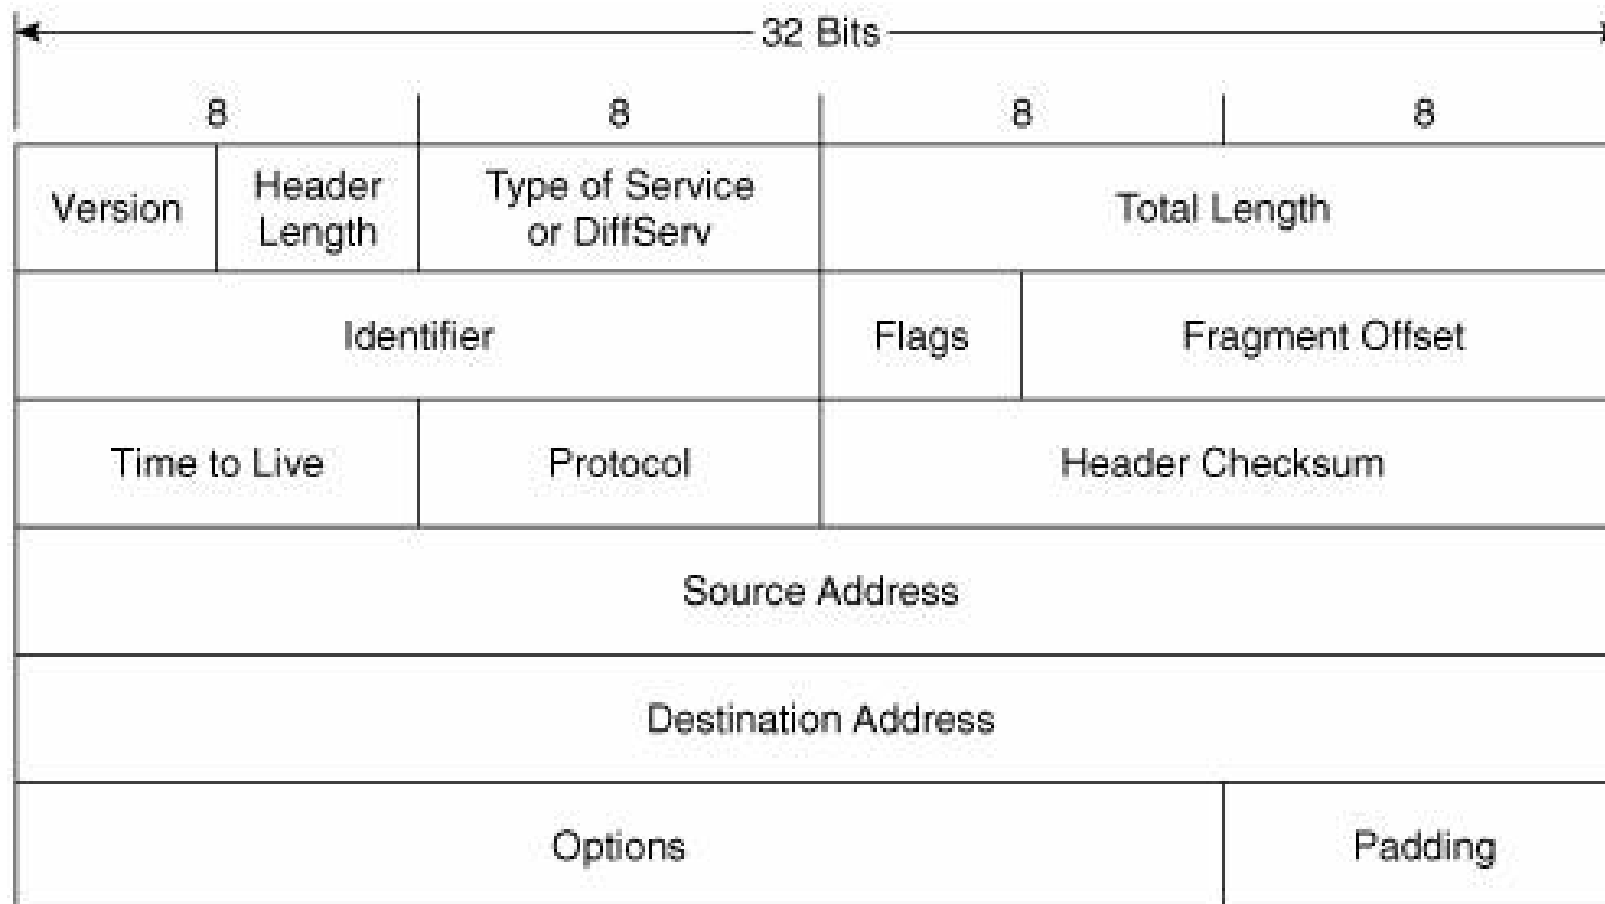
Fall 2019

**Dr. Seth James Nielson**

# IPV4 PROTOCOL

- Layer 3 Protocol
- Handles fragmentation and reassembly
  - Assumed that across multiple LANS, multiple MAC protocols
  - Each MAC protocol might have its own MTU
- Also, of course, includes the global IP address
  - Kind of global...

# IPV4 HEADER



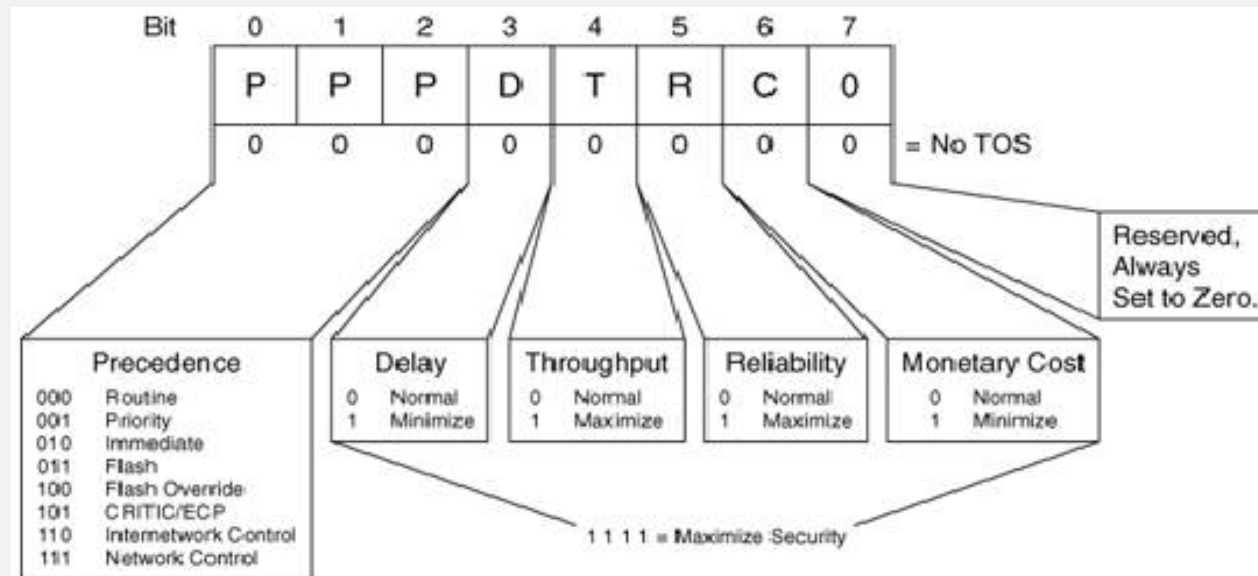
## IPV4 HEADER FIELDS

- Version – 0100 (binary 4)
- Header Length – Length of header in 4-byte increments
- Total Length – Size of header and data in bytes (max 65535)
- Identification – for recognizing fragments
- Flags
  - Reserved. Always 0
  - Don't fragment
  - More fragments

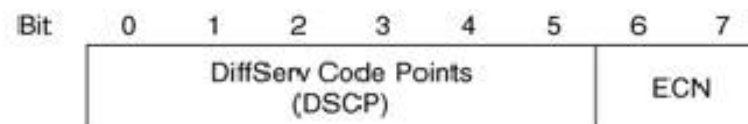
## MORE IPV4 HEADER FIELDS

- Fragment Offset – Offset in original datagram
- TTL – Counter to prevent infinite routing
- Protocol – Information about upper layer (17=UDP, 6=TCP)
- Header Checksum
  - Recomputed each hop (because of TTL changes)
  - Checksum field itself always presumed to be 0

# ORIGINAL TOS

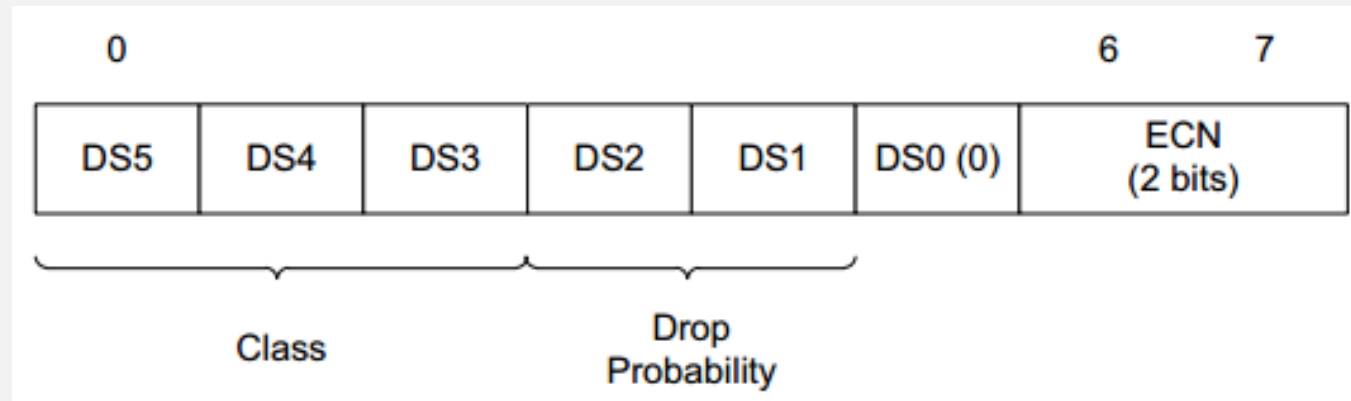


(a)



(b)

# DIFFERENTIATED SERVICES



# ECN

- Explicit Congestion Notification
- When enabled, indicates congestion without dropping
- Not all hardware/software supports ECN



# IPV4 ADDRESSES

- IPv4 Addresses are a.b.c.d where each is between 0-255
- In actuality, just a 32-bit number (“four octets”)
  - 192.0.2.235
  - 3221226219
  - 0xC00002EB (0xC0.0x00.0x02.0xEB)
- Private Networks:
  - 10.0.0.0
  - 172.16.0.0
  - 192.168.0.0

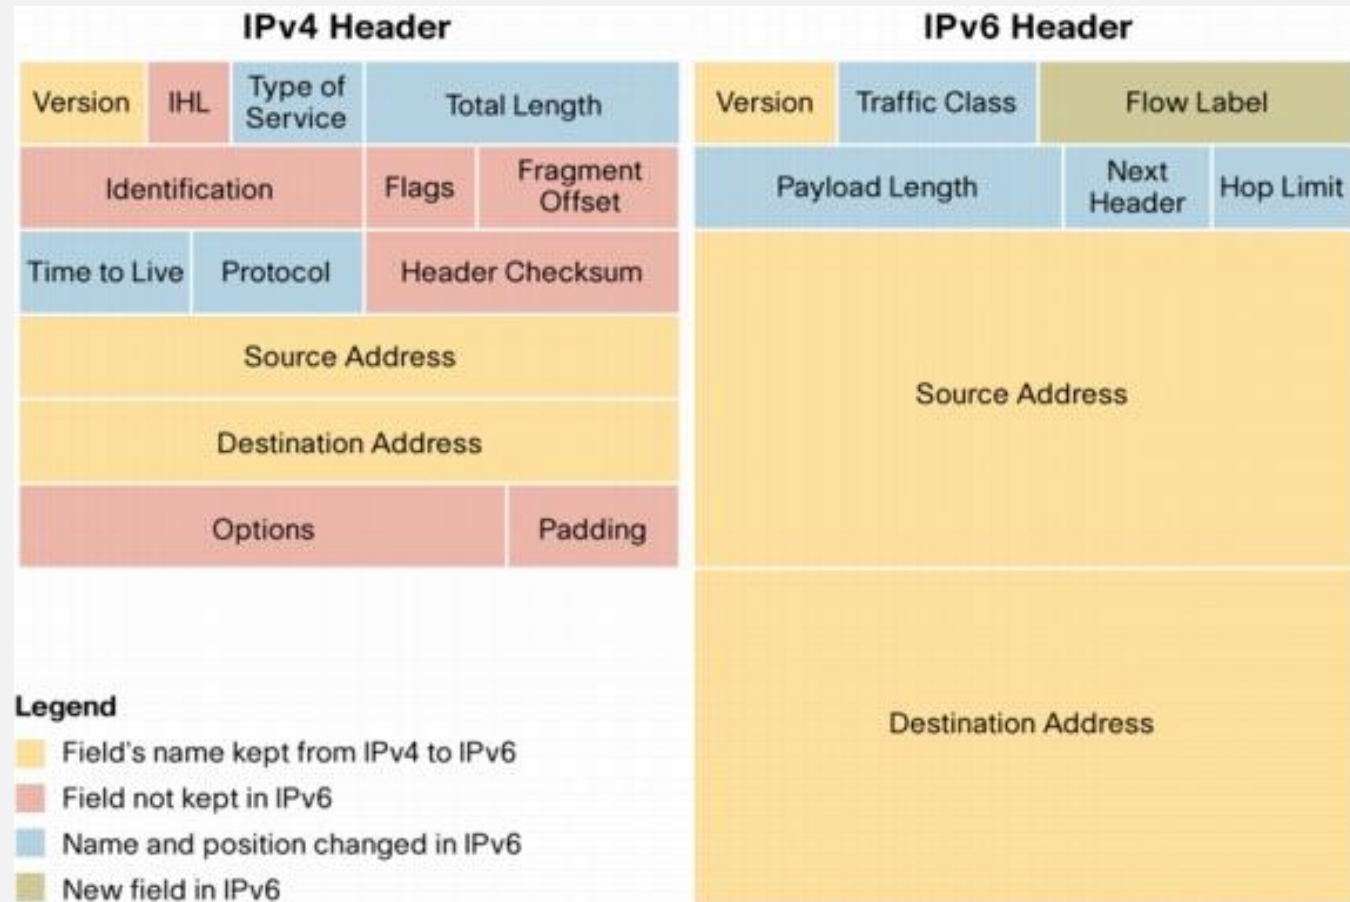
# FRAGMENTATION

- If packet size > MTU, and DNF is 0
  - Each fragment gets its own size
  - Each fragment except the last gets MF set to 1
  - Fragment offset is location in the original packet
  - Identification field is unique identifier of original datagram
- Reassembly
  - Use src, dst, protocol, and identification to identify fragments
  - Use offset to store data in reassembly buffer
  - Use MF = 0 to recognize end of reassembly

# FRAGMENTATION ISSUES

- IPv4 Fragmentation had security issues
  - IP Fragmentation Overlap (overwrite a fragment)
  - Buffer full (too many incomplete fragments)
  - Fragment overrun
  - (Note that most of these are DoS, but some evasion)

# IPV6

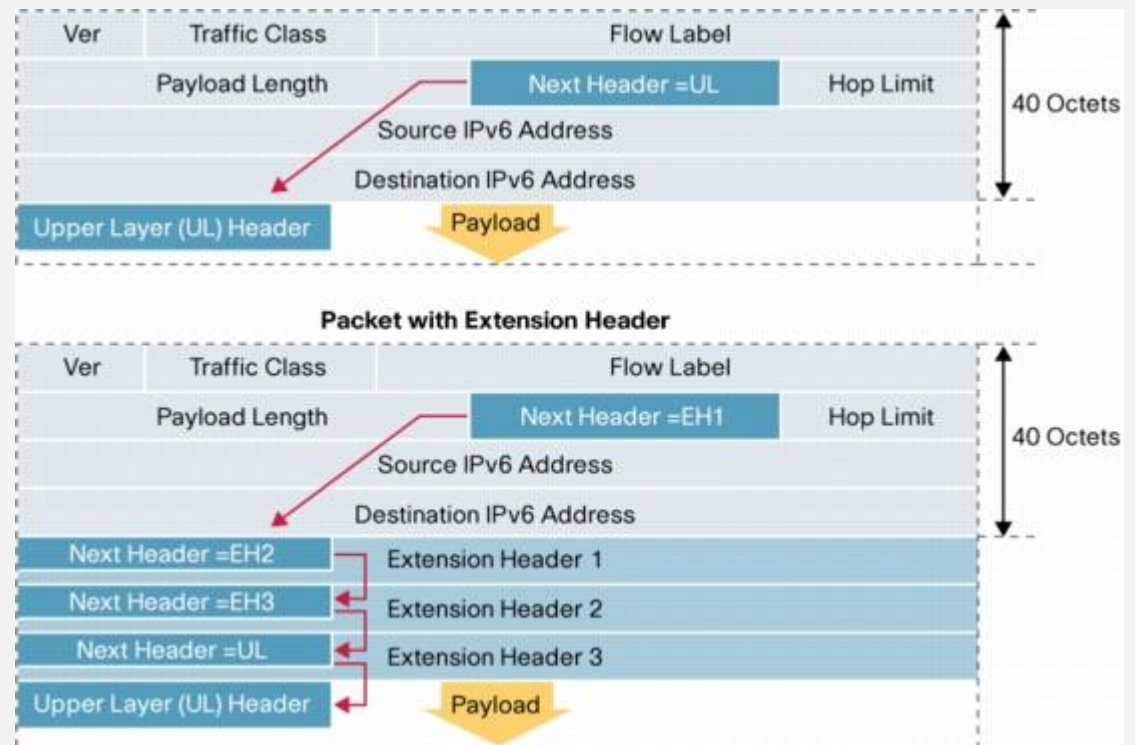


## IPV6 HEADER FIELDS

- Version = 0110 (binary 6)
- Traffic Class: Differentiated services plus ECN
- Flow label: Hint for multiple outbound paths
- Payload length: Includes extension headers, in bytes
- Next header: Type of next extension or transport header
- Hop limit: Decrement by 1, discard if 0

# EXTENSION HEADERS

- IPv6 is always 40 bytes for main header
- Can have additional headers:



# COMMON EXTENSION HEADERS

Order	Header Type	Next Header Code
1	Basic IPv6 Header	-
2	Hop-by-Hop Options	0
3	Destination Options (with Routing Options)	60
4	Routing Header	43
5	Fragment Header	44
6	Authentication Header	51
7	Encapsulation Security Payload Header	50
8	Destination Options	60
9	Mobility Header	135

# IPV6 FRAGMENTATION

- To deal with IPv4 Frag issues, ONLY SENDER can frag in IPv6
- Ergo, ***sender must know smallest MTU of path!***
- Path MTU Discovery (PMTUD)
  - If too big, send an ICMPv6 “packet too big” to sender
- Otherwise, max IPv6 packet size is 1,280 bytes.
- Uses a fragmentation extension header
  - Identification
  - MF, etc



# IPV6 FRAG PROBLEMS

- Studies from 2014-present indicate IPv6 fragmentation fails
- About one-third of IPv6 hosts could not receive frags
- Many are concluding that IPv6 fragmentation is deprecated
- Maximum IPv6 packet size between 1280 and 1350
- See,
  - <https://blog.apnic.net/2016/05/19/fragmenting-ipv6/>
  - <https://labs.apnic.net/?p=1033>

# TCP PROTOCOL

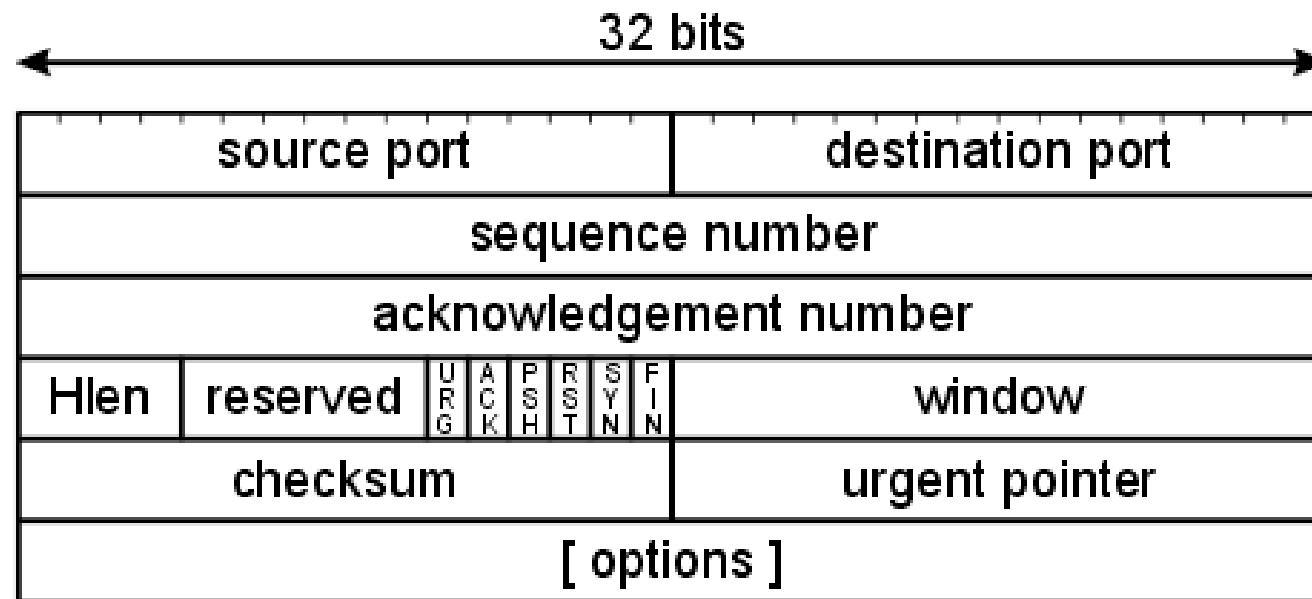
- Layer 4 Protocol
- Multiplexing (ports)
- Reliable Delivery (ack/resend)
- Congestion control
- Units called \*segments\*

# REVIEW

- Ethernet *frame*
- IP *packet*
- TCP *segment*

# TCP HEADER

## TCP header format



## TCP HEADER FIELDS

- Source Port, Destination Port for multiplexing
- Sequence Number of the first data byte
- Acknowledgement Number of next expected seq. no.
- Hlen number of 32-byte words in the TCP header
- Window number of bytes willing to receive
- Checksum over the header and data

## TCP FLAGS

- URG The URGENT POINTER field contains valid data
- ACK The acknowledgement number is valid
- PSH The receiver should pass this data to the application as soon as possible
- RST Reset the connection
- SYN Synchronize sequence numbers to initiate a connection.
- FIN Sender is finished sending data

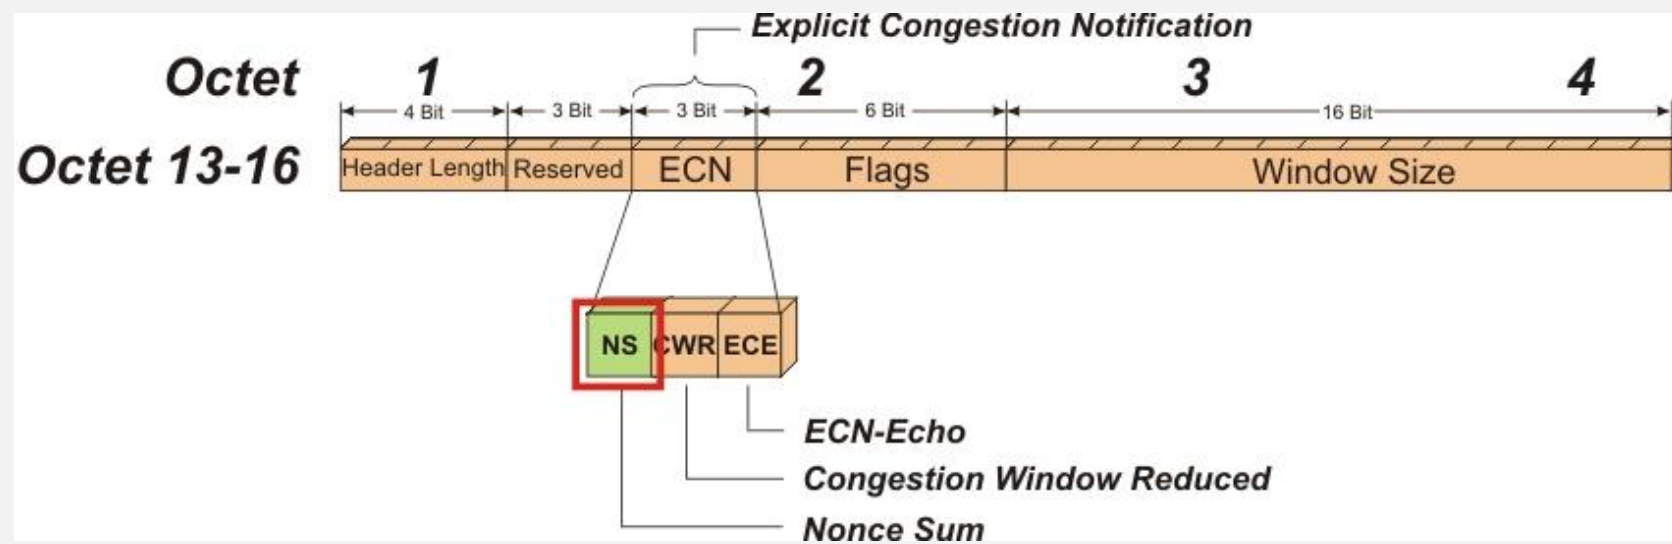


## TCP FLAGS (CONT)

- CWR – Acknowledges that congestion notification received
- ECN – indicates congestion notification via IP layer
  - (NOTE! Requires ECN capable IP layer!)
  - Sent until CWR received
- Only used if negotiated using TCP options during handshake



# OPTIONAL NS FLAG



## ONE-BIT NONCE

- NS is a parity bit used to catch changes to a packet
- Because it's only one bit, a cheater can guess it 50% right
- But, over repeated trials (frequent congestion) will get caught

# TCP OPTIONS

- Header can be “extended” with options
- Each option can have up to three fields:
  - Option Type (1 byte)
  - Option Length (1 byte)
  - Option data (variable)
- Examples include
  - Selective acknowledgement
  - ECN

# TCP SEQUENCE NUMBERS

- TCP header has a value for seq num and ack num every time
- SYN Sequence Number random between 0-4,294,967,295
- SYN Ack Num should be 0 (but any value should be ignored)
- SYN-ACK Seq Num also random
- SYN-ACK Ack Num is SYN Seq Num + 1
- (SYN-ACK) ACK Seq Num is SYN Seq Num + 1
- (SYN-ACK) ACK Ack Num is SYN-ACK Seq Num + 1

## DATA SEQUENCE NUMBERS

- Technically, there is only one packet type in TCP
- Flags simply indicate how the values can be used
- Sequence number is set every time
- But only increased by the length of the data
- (or increased by +1 for SYN and FIN)
- Ack field indicates that the ACK number is valid

# WIRESHARK TRACE

Time	192.168.1.2 174.143.213.18	Comment
0.000	(54841) → (80) SYN	Seq = 0 Ack = 94856056
0.047	(54841) ← (80) SYN, ACK	Seq = 0 Ack = 1
0.047	(54841) → (80) ACK	Seq = 1 Ack = 1
0.047	(54841) → (80) PSH, ACK - Len: 725	Seq = 1 Ack = 1
0.097	(54841) ← (80) ACK	Seq = 1 Ack = 726
0.100	(54841) ← (80) ACK - Len: 1448	Seq = 1 Ack = 726
0.100	(54841) → (80) ACK	Seq = 726 Ack = 1449
0.100	(54841) ← (80) ACK - Len: 1448	Seq = 1449 Ack = 726
0.100	(54841) → (80) ACK	Seq = 726 Ack = 2897
0.100	(54841) ← (80) ACK - Len: 1448	Seq = 2897 Ack = 726
0.100	(54841) → (80) ACK	Seq = 726 Ack = 4345
0.150	(54841) ← (80) ACK - Len: 1448	Seq = 4345 Ack = 726
0.150	(54841) → (80) ACK	Seq = 726 Ack = 5793
0.152	(54841) ← (80) ACK - Len: 1448	Seq = 5793 Ack = 726
0.152	(54841) → (80) ACK	Seq = 726 Ack = 7241
0.152	(54841) ← (80) ACK - Len: 1448	Seq = 7241 Ack = 726
0.152	(54841) → (80) ACK	Seq = 726 Ack = 8689

< ||| >

< ||| > ▼

Save As

Close

# TCP SHUTDOWN



# SECURITY

- Obviously, TCP is not designed with security in mind
  - No confidentiality!
  - No authentication!
  - No integrity (checksum is not cryptographic)
- No secure availability either!
  - End any connection with RST